

Time Series Introduction

Ordering

In general machine learning:

- no specific ordering of the data points
- all training points are potentially relevant for the prediction of a new point

In **time series**:



- chronological ordering of the points
- recent points are potentially more relevant than older points for the prediction of a new point

Patterns in time series

Generally, assume the time series can be decomposed into several elements, each of which we model separately:

- **Trend:** a long term increase or decrease in the data
- **Seasonality:** a cyclic pattern in the data (e.g.: days of the week, quarters of the year)
- **Noise:** a non deterministic element in the data

Seasonality vs. Trend





Hands-on session

`time_series_eda.ipynb`

>> Importing and visualising Time Series data

Resampling

Resampling involves changing the frequency of your time series observations.

Upsampling: increase the frequency of the samples, (e.g. from days to hours)

Downsampling: decrease the frequency of the samples (e.g. from days to weeks)

```
pandas.DataFrame.resample
```

Resampling

We use resampling because we have observations at the wrong *frequency*:

- They may be too granular or not granular enough!

Upsampling: this typically requires more care, as we are essentially interpolating between observations to *guess* what the measurements would have been in between.

For example: if we observe data hourly but need measurements every 30 minutes, then one approach could be to simply average the measurements before and after (linear interpolation).

```
pandas.DataFrame.resample
```

Resampling

We use resampling because we have observations at the wrong *frequency*:

- They may be too granular or not granular enough!

Downsampling: this typically *safer* than upsampling, because are aggregating data and reducing the granularity.

For example: we may observe hourly data but only be interested in daily measurements. Then one approach would be to compute the mean over all hourly measurements.

```
pandas.DataFrame.resample
```



CAMBRIDGE SPARK



Hands-on session

time_series_eda.ipynb

>> Resampling


>> Parsing custom date formats

Moving Windows

Big difference in time series: **not all data are equal!**

Moving windows involves applying a function on repeated fixed-width “slices” (window) of the data sliding along in the direction of the data

```
pandas.DataFrame.rolling
```



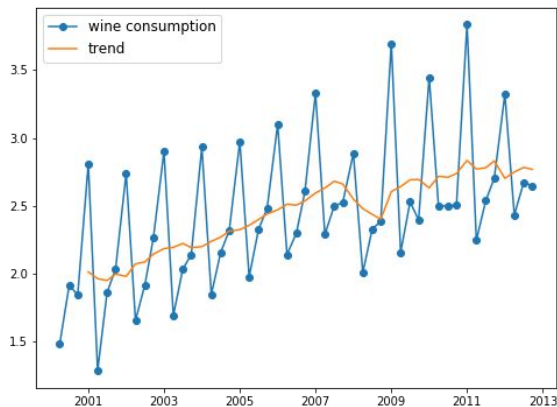
Window of width 4



CAMBRIDGE SPARK

Moving Average

- Better expose the trend of the data
- Creates a new series by calculating the average of fixed-width windows
- The window slides along the time series
- The same principle can be used for moving standard deviation, moving median etc





Hands-on session

`time_series_eda.ipynb`
>> Moving windows

Differencing

Differencing is a method of transforming a time series dataset.

Differencing is performed by subtracting the previous observation from the current observation. Here with a lag of 1:

$$\text{difference}(t) = \text{observation}(t) - \text{observation}(t-1)$$

Using differencing can help remove the trend and seasonality and expose the noise in the time series.

```
pandas.DataFrame.diff
```



Hands-on session

`time_series_eda.ipynb`
>> Differencing

Autocorrelation

Autocorrelation is the correlation (similarity) of a time series with a lagged version of itself.

Ex: take values `[1:10]` then values `[5:15]` how similar are these two sequences of values?

It helps expose the **seasonality** structure of the data.

```
pandas.DataFrame.autocorr
```



Hands-on session

`time_series_eda.ipynb`
>> Autocorrelation

Takeaways

- We often model time series data as consisting of a **trend** and **seasonal** component, together with some observational noise
- Often it may be helpful to resample the data in order to align it with our objectives. Upsampling requires more care than downsampling
- **Not all data are created equal!** We may discard past observations using methods such as sliding windows
- **Differencing** is helpful in removing some of the trend and seasonality in the data